

Peterson Guo

peterson.guo@uwaterloo.ca | petersonguo.com |  PetersonGuo |  PetersonGuo

EDUCATION

University of Waterloo

BASc. in Honours Electrical Engineering

Waterloo, ON

09/2023 - 04/2028

SKILLS

Languages: C/C++, Python, MLIR, Bash

Software: LLVM, CUDA, CuteDSL, Pytorch, Git

Others: GPU Kernels, Compilers, GPU Drivers, Linux, Kernel Debugging, LLMs

EXPERIENCE

Deep Learning Library Performance Software Engineer Intern

05/2026 – 09/2026

Nvidia

Santa Clara, CA

- Optimized mixed-precision MoE kernels in CuteDSL with Swap A/B, Split-K, 2CTA, scheduling, HBM prefetch, and PDL, improving low-latency inference by up to 46% vs. production baseline
- Added MXFP4 × MXFP8 support to grouped and block-scaled GEMMs, enabling mixed-precision MoE execution across CUTLASS, FlashInfer, and vLLM
- Improved Blackwell TGV GEMMs by fixing BF16 Split-K and NVFP4 TMA scale-factor loads, fusing bias, and tuning PDL launch counts for low-latency decode workloads

Model Performance Engineer Intern

01/2026 – 05/2026

Baseten

San Francisco, CA

- Optimized OpenAI Whisper inference on H100, achieving **5x** throughput and **60%** latency reduction via encoder batching, kernel fusion, and custom runtime scheduling on TRT-LLM
- Implemented LLM inference optimizations in TRT-LLM, integrating custom **CUDA** kernels using CUTLASS and CuteDSL, TurboQuant KV-cache quantization and RadixMLP scatter/gather kernels, improving memory efficiency and long-context/batched decode performance; presented kernel/runtime optimizations to FDEs and led technical reviews
- Re-architected speaker diarization pipeline, delivering **15x** throughput and **80%** latency reduction through parallelization and model-level optimizations for a **multi-million** dollar enterprise deployment
- Built a privacy-preserving LLM request hashing + replay system for deterministic workload re-simulation, enabling performance regression testing and capacity planning

Compiler Engineer Intern

05/2025 – 12/2025

Nvidia

Toronto, ON

- Built and owned an internal **Python-C++** profiler for CUTLASS latency decomposition with up to **21x** lower overhead than torch.profiler and adopted by **15+** systems developers
- Reduced end-to-end **JIT** compilation time by **20%** by designing a deterministic Merkle-tree hashing system **5x** faster than SHA256 for multi-GB binaries, improving kernel generation **latency** and iteration speed
- Reduced kernel launch overhead by **60%** by eliminating **Python→C++→CUDA** dispatch hotspots, achieving performance within **10%** of tvn-ffi and over **2x** faster than Triton
- Core contributor to an **MLIR**-based CUDA dialect, designing and lowering **20+** ops to enable Python DSL access to graphs, streams, and async memory operations

Software Engineer Intern

09/2024 – 12/2024

AMD

Markham, ON

- Delivered AMD's most stable GPU software release, by resolving **25+** kernel-level crashes, hangs, and performance regressions, using crash dumps, **firmware** logs, and **register-level** analysis
- Shaped display-pipeline research, shown by my **ML upscaling** PoC being adopted and later expanded into an internal technical conference paper by a successor, by prototyping a lightweight **PyTorch/ROCm** super sampling model
- Strengthened display-pipeline stability on next-gen AMD GPUs/APUs by building **C/C++ kernel** drivers, ensuring robust framebuffer→display behavior for Ryzen AI and Radeon hardware
- Contributed upstream fixes to AMD's open-source **Linux** display driver, resolving defects in color calibration, frame synchronization, DSC decoding, and visual corruption

Security Developer Co-op

01/2024 – 05/2024

eSentire

Waterloo, ON

- Performed **LLM security research** on production backends, analyzing attack surfaces including prompt injection, data exfiltration, and privilege escalation in deployed pipelines
- Cut **ingestion latency** by **50%+** by **optimizing** JSON parsing and Snowflake execution paths, reducing pipeline stalls in **real-time analytics**
- Improved log-processing throughput by **4x** through **algorithmic optimizations** in **Python's** data pipeline

PROJECTS

Chess Bot | PyTorch, Python, C++, CUDA, Alpha Beta Pruning, NNUE

- Built a **GPU-accelerated** chess engine using custom **CUDA** kernels for evaluation/search

InvestIQ | Python, C++, IBKR API, CUDA

- Built a low-latency, event driven market data pipeline aggregating price, event, and sentiment feeds, for backtesting and feature generation with pairs trading and signal research